

Satyam Kumar

+91 9693342101 satyamkr1354@gmail.com linkedin.com/in/satyam1354 github.com/satyamchy
satyam-portfolio.vercel.app

Summary

Software Engineer with 1+ year of experience building and shipping AI-powered products and full-stack applications in production. At Xornor Technology, I engineer LLM-driven workflows, REST APIs, and scalable backend services across real user-facing products. Earlier, at Innova Solutions, I built automated data pipelines processing enterprise-scale healthcare data for a Fortune 25 US client. Core stack: Python · React · Node.js · FastAPI · LangChain · PostgreSQL · AWS. Seeking a Software Engineer or AI Engineer role where I can build intelligent, scalable systems from the ground up.

Technical Skills

Languages: Python, JavaScript (ES6+), TypeScript, SQL, Java, C/C++
Web & APIs: React.js, Node.js, FastAPI, REST APIs, PostgreSQL, Redis, JWT, System Design
AI/LLMs: LangChain, RAG, Fine-tuning, Vector Search, Agentic AI, PyTorch
Cloud & DevOps: AWS (EC2, S3), Docker, GitHub CI/CD, Kubernetes
Data Engineering: Snowflake, PySpark, ETL Pipelines, Power BI, Power Automate
Tools: Git, Postman, Linux, VS Code, Agile / Scrum

Experience

- | | |
|--|---|
| Software Engineer — AI & Full Stack
<i>Xornor Technologies</i> | May 2026 – Present
<i>Mohali</i> |
| <ul style="list-style-type: none">Built and deployed an AI-powered Jira analytics platform end-to-end — Jira REST API ingestion, LangChain + OpenAI pipeline for worklog classification and anomaly detection, and a React dashboard surfacing productivity and utilisation insights to engineering managers across 50+ projects.Delivered within a 2-week sprint — architected idempotent cron-scheduled ETL pipelines for continuous Jira data synchronization, implemented deduplication and schema validation layers ensuring zero data loss across daily ingestion cycles, and eliminated 90% of manual reporting overhead.Built KAI-LMS, a production multi-role Learning Management System with an AI-driven rubric-based grading module; the LLM evaluation layer assesses open-ended student responses against structured rubric criteria, providing consistent detailed feedback at scale and cutting instructor review time by an estimated 10×.Currently extending KAI-LMS as part of a 4-engineer team — engineering a real-time in-app messaging system, integrating LLM-based automated evaluation for assignments and quizzes, and resolving production bugs across the platform to improve system stability and user experience. | |
| Software Engineer Intern (Data Pipelines)
<i>Innova Solutions</i> | May 2025 – March 2026
<i>Hyderabad</i> |
| <ul style="list-style-type: none">Built an automated Power Query pipeline to ingest, normalise, and consolidate 1,800+ monthly CSV files across 60+ heterogeneous report schemas into a single unified analytics dataset, eliminating all manual report merging and enabling consistent month-over-month analytics at scale.Replaced legacy SAS data extraction to optimised Oracle SQL queries and built an end-to-end workflow (Oracle DB → Power Automate → SharePoint → Power BI), enabling daily & monthly refresh dashboards tracking Medicaid and Medicare insurance claim errors across public and private payers.Built a derived Error Analytics Dataset by engineering computed columns — Key ID, Entry Date, Claim Age, and SLA metrics — enabling compliance trend visualisation and SLA breach tracking across monthly reporting cycles for a client processing millions of healthcare claims annually. | |

Projects

- AI Portfolio Assistant** | *LangChain · RAG · ChromaDB · React · Node.js · Express* Apr 2026
- Built a portfolio chatbot with a RAG pipeline — stored portfolio data in ChromaDB, retrieved relevant chunks on each query, and passed them as context to the LLM for source-grounded answers; cuts prompt size compared to sending full data on every call.
 - Deployed full-stack with React frontend and Node.js/Express backend; handles real-time queries with conversational follow-up support.
- AI-Powered News Aggregator** | *FastAPI · LangChain · Groq (Llama 3) · React* March 2026
- Built a full-stack AI news platform where a LangChain React agent autonomously searches the web via DuckDuckGo, synthesising real-time snippets from 10–20 sources into concise cited summaries in under 3 seconds; engineered category-aware prompt templates for structured JSON output.
 - Implemented an in-memory caching layer (TTL 30 min) reducing redundant Groq API calls by 85%; deployed via GitHub Actions CI/CD — React on Vercel, FastAPI on Render — at \$0/month infrastructure cost.
- Full Stack Developer — QuantFit** Feb 2025
- Developed a fitness web platform with browser-based PoseNet pose detection for real-time rep counting and movement accuracy tracking; integrated BMI calculation, calorie monitoring, and personalised fitness insights with a React dashboard and Node.js/Express backend for secure user data persistence.

Education

Chitkara University Baddi, HP
Bachelor of Engineering in Computer Science - 9.54 CGPA *Aug 2022 – May 2026*

Key Achievements & Certifications

- IBM AI Fundamentals · AI for Scientific Research · Software Product Management — Coursera / IBM
- Delivered production-grade data pipelines for Centene (Fortune 25 US healthcare) as an entry-level engineer, processing Medicaid and Medicare claim data at enterprise scale.
- Applied Snowflake and Informatica Intelligent Cloud Services (IICS) in practical data integration tasks during internship at Innova Solutions.
- Graduated 9.54 CGPA from Chitkara University while working full-time as a Software Engineer Intern at Innova Solutions.